

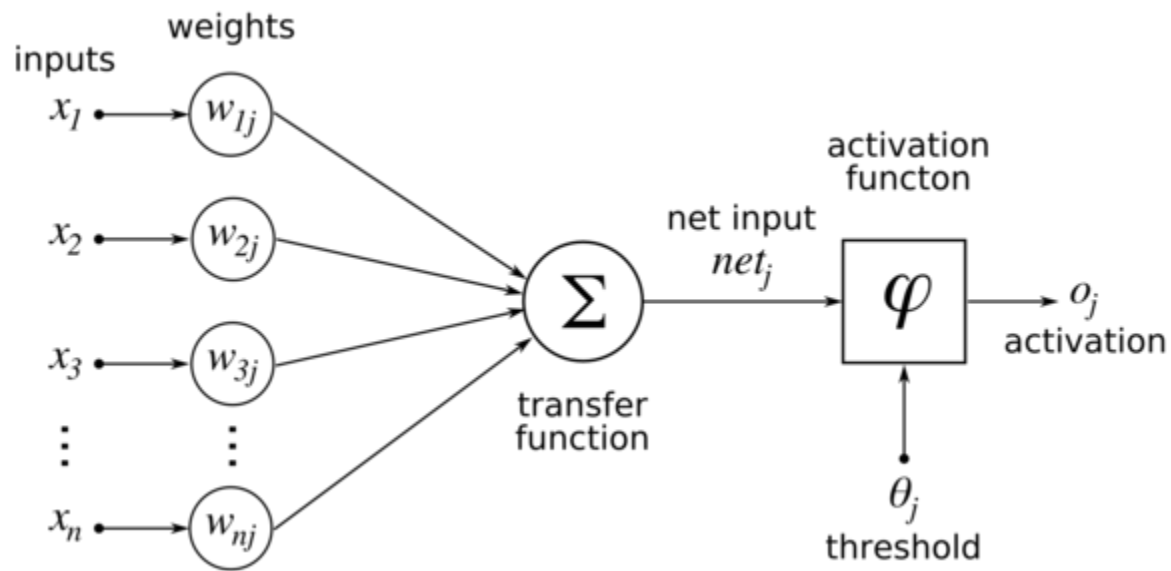
Neuronale Netze - Eine kurze Einführung mit Implementierungen in Python

Philipp Hanemann, Martin Czygan

The origin - a linear classifier



$$H_i = \langle x, w \rangle = \sum_j x_j w_j + \theta$$



Activation function can vary e.g.:

- step function

\$\$

$o_j =$

$\begin{cases} 1, & \text{if } net_j \geq 0 \\ 0, & \text{else} \end{cases}$

\end{cases}

\end{cases}

\end{cases}

\$\$

How to obtain the weights?

The objective is a good model fit.

- trial an error \rightarrow inefficient
- optimization $\rightarrow \min_w \text{Cost}(t, w)$

with:

t : target value

w : weight vector

e.g. squared error as in linear regression

\rightarrow optimization theory

- one efficient way for solving the problem is the use of backpropagation (error is "propagated" backwards through the grid)

Representing Boolean Algebra as Classifiers

x_1	x_2	AND	OR	XOR
0	0	0	0	0
0	1	0	1	1
1	0	0	1	1
1	1	1	1	0

AND is linearly seperable



One possible AND perceptron



OR(/NOR) is linearly seperable



One possible AND perceptron



XOR is not linearly separable



XOR can be represented by a combination of two mappings

XOR = NOR (AND, NOR)

point	x_1	x_2	$(x_1 \text{ AND } x_2)$	$(x_1 \text{ NOR } x_2)$	NOR	XOR
a	0	0	0	0	1	0
b	0	1	0	1	0	1
c	1	0	0	1	0	1
d	1	1	1	0	0	0

**The extra mapping can be
visualized**



One possible XOR Net (#1)

The ones are fixed input (bias) units



One alternative XOR Net (#2)

The number within the perceptron represents the inherent bias unit/or a translational shift when the unit jumps.



Two nets with the same result - why care?

	Net #1	Net #2
# weights	9	5

- Net #1 has more free variables
- Net #1 has a higher dimensional weight space (\mathbb{R}^9 vs. \mathbb{R}^5)
- Net #2 has less degrees of freedom and should generalize better.

Why is that?

- This architecture of the net has a direct effect on the optimization problem and the search space.

The MNIST Dataset for benchmarking



Playing with MNIST and scikit-learn

```
from sklearn.neural_network import MLPClassifier
from sklearn.datasets import fetch_mldata

MNIST = fetch_mldata("MNIST original")
split = 60000 # number of training examples
X, y = MNIST.data / MNIST.data.max(), MNIST.target
X_train, X_test = X[:split], X[split:]
y_train, y_test = y[:split], y[split:]
mlp = MLPClassifier(hidden_layer_sizes=(n_units, n_layers),
                    max_iter=n_iterations, alpha=1e-4, solver=solver,
                    verbose=10, tol=1e-4, random_state=1,
                    learning_rate_init=alpha)
mlp.fit(X_train, y_train)
score = mlp.score(X_test, y_test)
```

Evaluating the parameter space

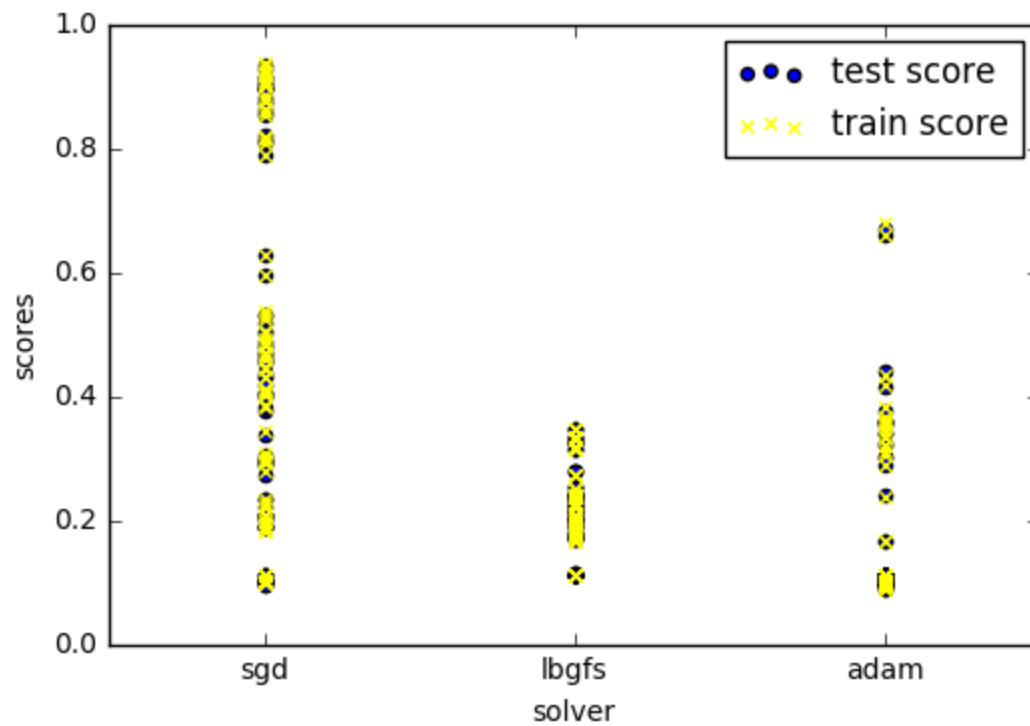
cartesian product of:

```
hidden_units_per_layer = [2, 4, 6, 10, 15]
hidden_layers = [1, 2, 3]
learning_rate = [0.1, 0.2, 0.3]
solver = ['lbfgs', 'sgd', 'adam']
max_iter = [5, 10]
```


\$\rightarrow\$ 270 runs

```
base_config = {'hidden_units_per_layer': 10.0,
               'iterations': 10.0,
               'hidden_layers': 1.0,
               'solver': 'sgd',
               'learning_rate': 0.1}
```

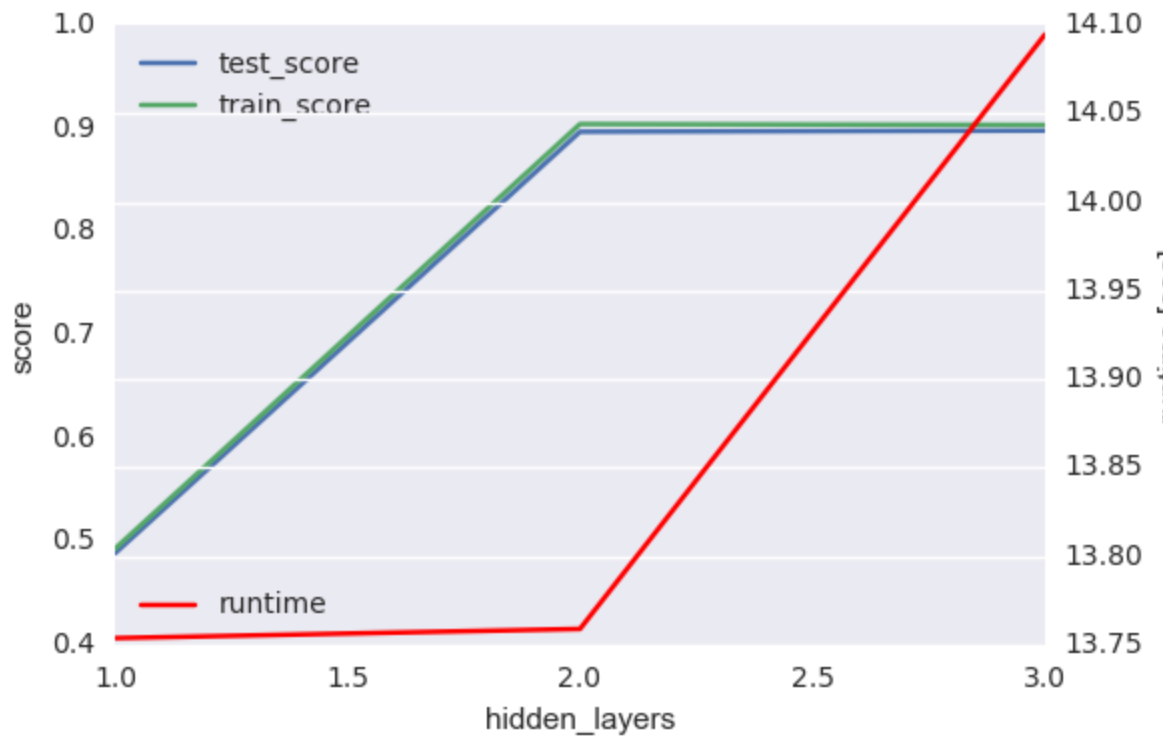
Influence of the solver




Number of hidden units per layer




Number of hidden layers



Number of iterations



Influence of the learning rate



Some Code

Roadmap:

- [perceptron.py](#)
- [randomweights.py](#)
- [pocket.py](#)
- [xorish.py](#)
- [basicnn.py](#)
- [mnistimages.py](#)
- [hellosklearn.py](#)
- [sknngrid.py](#)
- [hellotf.py](#)
- [hellokeras.py](#)

perceptron.py

A simple perceptron plus lots of boilerplate for gif.

- simple update rule
- relatively fast, given the weight space is infinite
- works on separable data

randomweights.py

